

INSIGHTS PRESENTATION

Claude Code Usage Analytics Platform

From raw telemetry to interactive insights

How we reached these results

We start from raw telemetry, normalize it into a warehouse, then read the same aggregates in Streamlit that power the findings on the following slides.

The path


1. **Ingest:** CSV (employees) + JSONL (events) → parsing & validation (Pydantic) → **PostgreSQL** tables.
2. **Derive:** Events are rolled into **sessions** (duration, counts, tokens) so we can talk about “work units,” not only raw lines.
3. **Consume:** The **gateway** exposes `/metrics` (and related routes); the dashboard applies the **same filters** everywhere (dates, practice, location) so every chart is comparable.

What we measure — and why

Tokens	Proxy for cost and model load; sums align with how vendors bill.
Events	Granular activity (prompts, tool use, API calls); explains intensity beyond one session.
Sessions	Bounded “visits”; pairs with duration buckets to describe how long people stay in a session.
Hour / day (UTC)	When usage happens—helps separate work-hours from off-hours.
Practice / level / location	Org dimensions from employee metadata; shows where demand concentrates.
Event type & model	What happened (e.g. tool vs API) and which model —links behavior to routing and cost.

When • Usage follows working hours (UTC)

Token volume is elevated during daytime UTC and lower overnight — consistent with work-aligned usage.




Streamlit · Overview · Tokens by hour of day (UTC)

Pattern: Higher activity through the workday band in UTC.

Read: Usage concentrates in periods that resemble on-the-job work, not evenly across 24 hours.


Who • Practice & seniority concentration

Practice first, then seniority (same filters).



Tokens by practice

By practice: ML and Frontend are highest; Data and Backend are mid-pack; Platform is lowest here—compare to team size before inferring “adoption.”




Tokens by level

By seniority: Most tokens sit in L3–L6. L1–L2 and L7+ are lower—often fewer people and/or lighter usage per person.


How • Short sessions & model mix

Session shape first, then which models show up in events (same filters).



Session duration distribution

Sessions: Most sessions sit in the shortest duration buckets—many brief sessions rather than a few long ones. That fits task-sized use (quick questions, edits, runs) and can still add up to high event volume across the day.



Model usage (events)

Models: Claude Haiku 4–5 leads in event counts—consistent with **cost/speed** tradeoffs in the mix, not a claim about “best” quality.

IN SHORT

If you read the slides above as one thread—same data, same filters—the picture is coherent: usage **isn’t scattered across the clock**; it **ramps up through daytime UTC**, which reads like people reaching for the assistant during real work blocks, not at random.

Under the surface, **who** shows up is uneven: **mid-level** engineers carry much of the token load, and **practice** bars show where the org pushes hardest. **How** that work happens skews toward **short sessions**—many quick bursts—and a **Haiku-heavy** event mix, which fits **task-sized** help and **cost/speed** tradeoffs. That’s the story of *when*, *who*, and *how* this assistant lands in the day.

Synthetic data; patterns are illustrative. UTC vs local time; correlation ≠ causation.